Building M&A sourcing on a bootstrapped budget

You can build all four sourcing features for roughly $60–80/month — well under the $225 ceiling — by combining free government APIs, Google Places' generous free tier, Claude API for intelligent enrichment, and a curated-first approach to professional and broker data. The critical insight: the most valuable data sources for SMB M&A (Census Bureau, SEC EDGAR, SBA loan data, SBIC directory) are completely free and public domain, while the expensive commercial databases (PitchBook, Grata, Crunchbase) can be approximated at launch through AI-enriched public data and user-contributed content.

This report covers the specific data sources, APIs, legal constraints, architecture patterns, and a week-by-week build plan for each of the four features: acquisition target sourcing, buyer/investor matching, professional services matching, and a broker collaboration hub.

Acquisition target sourcing runs on three free data layers

The buyer-side sourcing feature — helping buyers find businesses to acquire — requires combining multiple data sources because no single free source provides both business identification and acquisition-relevant signals.

Layer 1: Census Bureau for market intelligence (free). The County Business Patterns and ZIP Code Business Patterns APIs provide establishment counts, employment, and payroll data by **6-digit NAICS code and geography**. U.S. Census Bureau A query like `api.census.gov/data/2023/cbp?get=ESTAB,EMP,PAYANN&for=county:*&in=state:04&NAICS2017=238220` returns every plumbing contractor establishment in Arizona by county. This powers the market-density layer: "There are 847 HVAC establishments in Maricopa County employing 12,400 people." The API is free with a key (register at api.census.gov), supports 500+ queries/day, and returns JSON. U.S. Census Bureau Data is current to 2023 reference year. U.S. Census Bureau This doesn't identify specific businesses, but it frames the opportunity and validates a buyer's thesis.

Layer 2: Google Places for real-time business discovery ($0–50/month). After Google's March 2025 pricing overhaul, Nicola Lazzari the free tier works like this: Nicola Lazzari **Text Search Essentials (IDs Only) is unlimited and free**, google returning Place IDs for queries like "HVAC companies in Phoenix." Place Details Essentials then costs $5/1,000 requests with **10,000 free per month**, Nicola Lazzari google returning name, address, phone, and hours. Place Details Pro (adds ratings, website, reviews) costs $17/1,000 with 5,000 free monthly. google At 50–200 users running 5 searches each, you'll consume 250–1,000 text searches and perhaps 3,000 detail lookups — comfortably within free caps if you use field masks carefully. Requesting even one Pro-tier field upgrades the entire request's billing tier, Google so the Node.js implementation must enforce strict field masks.

The critical Google Places ToS constraint: you **cannot store Place Details in your database** to build a business listings dataset. Google explicitly prohibits caching content beyond Place IDs (which can be stored indefinitely). Google The compliant architecture is: store Place IDs in PostgreSQL, re-fetch details on demand, display with Google attribution. This means Google Places powers live search, not a persistent database. Place IDs should be refreshed every 12 months. Google

Layer 3: SBA loan data for acquisition signals (free, public domain). The SBA publishes complete 7(a) and 504 loan data Data.gov Sba at data.sba.gov going back to 1991, updated quarterly. Fields include borrower name, address, NAICS code, loan amount, lender, and approval date. SBA 7(a) loans are the primary financing vehicle for small business acquisitions in the U.S. — roughly **40% of SBIC financing flows to acquisition capital**. ZMC By filtering for loans in specific NAICS codes and geographies from 5–10 years ago, you can identify businesses that were acquisition-financed and whose owners may now be approaching exit. This data is public domain with zero licensing restrictions.

What the competitors actually do. Grata (acquired by Datasite in 2025) Grata built its 19M+ company database by scraping company websites with NLP, reading structure, language patterns, and commercial signals to classify what businesses actually do. Grata SourceScrub aggregates 290,000+ sources including trade show attendees, conference lists, product reviews, and portfolio company lists, Sourcescrub verified by 500+ industry specialists. Grata Axial takes a fundamentally different approach — it's a private network where sell-side advisors list opportunities and buy-side members describe criteria, with an algorithm matching confidentially Mademarket +3 across **10,000+ deals annually**. The Silicon Review +2 Baton ("Zillow for small businesses") uses public government data to generate 4M+ business estimates, then gets sellers to connect QuickBooks for financial verification. Built In NYC The lesson: **start with public data, then build proprietary data through user participation** — exactly what a bootstrapped platform should do.

Legal position is strong for public business data

Aggregating public business entity data and presenting it to users ("here are 10 HVAC companies in Phoenix that might be acquisition targets") is **generally legal and low-risk**. CCPA defines "personal information" as data linked to natural persons CA — business entity information (company name, address, industry, formation date) is not personal information. CCPA explicitly exempts "publicly available" information Clarip from federal, state, or local government records. Census data, SBA FOIA data, and Secretary of State filings all qualify. A bootstrapped startup below CCPA's thresholds ($25M revenue, 50,000+ consumers' data, or 50%+ revenue from selling personal information) likely isn't subject to CCPA regardless. Ketch

Scraping BizBuySell, LoopNet, or BusinessBroker.net is high-risk and not recommended. CoStar Group (which owns BizBuySell, BizQuest, and LoopNet) is notoriously litigious about scraping. While the *hiQ v. LinkedIn* precedent suggests scraping public data doesn't violate the

CFAA, (California Lawyers Associa...) courts have enforced website Terms of Service as breach-of-contract claims. (Morgan Lewis) For a bootstrapped startup, even winning a legal battle would be financially devastating. The practical boundary: **use government data and API-accessed data freely; don't scrape commercial listing sites**.

Buyer matching starts with four free databases

When a seller lists on the platform, the buyer/investor matching feature identifies PE firms, search funds, family offices, and individual buyers whose criteria fit. Four free data sources provide a credible seed database.

The SBIC Directory is the single highest-value free source. The SBA publishes a downloadable CSV of all ~371 licensed Small Business Investment Companies at sba.gov/funding-programs/investment-capital/sbic-directory. (SBA) Fields include company name, city/state, managing entity, fund size, average investment size, investment strategy (equity/debt/hybrid), fund style (venture/buyout), whether they're making new investments, and investor relations contact information. (ZMC) SBICs focus on businesses with net worth under $18M and average after-tax income under $6M (ZMC) — precisely the SMB acquisition sweet spot.

SEC EDGAR identifies active acquirers for free. The EDGAR Full-Text Search system (efts.sec.gov) enables free, unauthenticated boolean searches across all filings since 2001. (SEC) Query for 8-K filings containing "completed acquisition" filtered to Item 2.01 to identify companies that actively acquire businesses. The Submissions API (data.sec.gov/submissions/CIK{cik}.json) returns complete filing histories in JSON. (SEC) Form ADV filings reveal PE firms and family offices: the Investment Adviser Public Disclosure database contains **15,000+ registered advisers** with assets under management, client types, and strategies. Rate limit is 10 requests/second with just a User-Agent header — no key required. (Sec)

Search fund directories provide 1,700+ targeted buyers. The search fund ecosystem has exploded: **94 core search funds launched in 2023 alone** (Stanford 2024 Study), (Stanford Graduate School ...) with a typical acquisition profile of $1M–$10M EBITDA, median purchase price $14.4M. Axial's public directory lists 1,726 search funds with firm names, descriptions, and investment criteria. (Axial) (Axial) Searchfunder.com maintains community profiles and an annual formation report. (Searchfunder) Curated lists from Leland and SMEVentures name the top investors and accelerators (Pacific Lake Partners, Anacapa Partners, Relay Investments, Search Fund Partners, (Leland) among others). This data can be manually compiled once and enriched via AI website analysis.

A free PE/VC dataset seeds broader coverage. Gigasheet publishes a downloadable CSV of 23,916 VC/PE businesses globally with country, industry, founding year, and website. (Gigasheet) While it lacks investment criteria, combining it with AI-driven website scraping (reading each firm's "Investment Criteria" page) produces structured data on industry focus, deal size

preferences, and geography. This enrichment pipeline — fetch website, extract criteria via Claude, store structured output — costs roughly $0.45 per 100 profiles using Claude Haiku.

The matching algorithm mirrors Axial's approach

Axial's matching (Privsource) (revealed through their public GitHub coding challenge) uses **multi-dimensional overlap scoring**: each buyer and seller has associated geography/industry pairs, with edges weighted by the number of shared criteria. (GitHub) For smbx.ai's MVP, the implementation is straightforward:

1. Parse seller's NAICS code → map to industry categories
2. Parse revenue/EBITDA → map to deal size tiers ($500K–$2M, $2M–$5M, $5M–$10M, $10M+)
3. Parse geography → map to regions
4. Score each buyer: `match_score = (industry_overlap × 3) + (size_overlap × 2) + (geography_overlap × 1)`
5. Boost by buyer type relevance (PE for $5M+ EBITDA; search funds for $1–5M; individuals for <$1M)
6. Use Claude to generate natural-language match explanations ("This PE firm acquired 3 similar businesses in your industry over the past 2 years")

The differentiator versus Axial isn't database size — it's **AI-generated match reasoning, no success fees** (Axial charges Lehman-scale fees), (Privsource) and integration with the seller's exit-readiness journey. A seed database of 5,000–25,000 buyer profiles from free sources delivers roughly **60–70% match relevance** compared to Axial's estimated 80–85%, which is sufficient for launch.

Professional matching works best as curated-first, then automated

The professional services feature connects users to M&A attorneys, transaction CPAs, business appraisers, and SBA lenders. The MVP should be **manually curated** (5–10 verified providers per category per major metro), with automated discovery layered on in month two.

Google Places API can find professionals but can't verify specialization. Text Search handles queries like "mergers acquisitions attorney Austin TX" with surprisingly relevant results because Google indexes business descriptions and website content. (Google) However, the API's place type filters are broad — `lawyer` and `accounting` exist, but "M&A attorney" or "business appraiser" do not. (Google) The AI enrichment layer becomes essential: for each Google result with a website, Claude scans the homepage and about page for M&A-specific keywords ("mergers and acquisitions," "business sale," "due diligence," "SBA 7(a)," "business valuation") and assigns a

relevance score. This converts a generic attorney search into a specialized M&A professional discovery tool at minimal cost.

SBA lender data provides authoritative, free lender rankings. The SBA publishes monthly 7(a) Lender Activity Reports ranking every participating lender by loan count and dollar volume, Sba segmented by state. In FY2025, over 23,000 lenders completed at least one 7(a) loan, but only **1.6% completed 10+ loans**. Bankrate The top 25 lenders (Live Oak Bank, Huntington, U.S. Bank, TD Bank, Wells Fargo, Celtic Bank) account for 75%+ of total volume. Sbalenders Importing this data into PostgreSQL and matching by user location and loan size produces an authoritative lender recommendation at zero cost.

Professional directories exist but lack APIs. The NACVA (National Association of Certified Valuators and Analysts) has 7,000+ members SCORE searchable by state, NAICS code, and specialty — the most comprehensive business appraiser directory. The ASA (American Society of Appraisers) directory supports search by appraisal discipline Appraisers including Business Valuation specifically. IBBA's public directory lists business brokers by state and industry. Businessbrokeragepress Every state bar association has an online attorney search. Texas Board of Legal Speci... None offer public APIs, but Avvo's model proves the concept: Avvo scraped all 50 state bar associations to auto-create profiles for ~1.3M attorneys without consent, then enriched with court filings and licensing data. Wikipedia For smbx.ai, the practical approach is linking to these directories rather than replicating them, while building a curated layer of verified M&A specialists.

The Lendio model is particularly instructive: borrower fills out one application, proprietary algorithm matches against 75+ lender criteria, Vizologi funding manager contacts with matched offers. Lendio has facilitated **$16B+ in funding for 400,000+ businesses** Bestbusinessloansonline and offers an embeddable marketplace widget. Lendio This "single input → multi-provider match" pattern maps directly to smbx.ai's professional matching feature.

The broker hub fills a gap no one else has solved

Unlike real estate — which has true MLS systems with standardized data (RESO Web API, Medium OData Matrix v4.0), co-brokering protocols, and shared commissions — business brokerage has **no unified MLS**. Brokers list on BizBuySell ($40–60/month), Synergybb BusinessBroker.net ($40/month), Businessbrokeragepress and DealStream independently. International Business Bro... The IBBA doesn't operate a central deal exchange. International Business Bro... Co-brokering is informal: a buyer's broker contacts the listing broker directly, and commission splits (typically 50/50 on the standard 10–15% fee) are negotiated ad hoc. Viabeacon This fragmentation is the opportunity.

Three-stage information disclosure protects confidentiality

Confidentiality is the paramount concern in business sales — a leak can cause employees to quit, customers to defect, and competitors to poach. Michael Shea The broker hub must enforce a strict three-stage model:

- **Stage 1 (Blind Profile/Teaser):** Industry category, broad geographic region, revenue range, EBITDA range, employee count range, and a sanitized business description. No identifying information whatsoever — no name, exact address, specific products, or unique characteristics. OffDeal Visible to all authenticated brokers.
- **Stage 2 (CIM Access, after NDA):** Business name, exact location, 3 years of detailed financials, customer concentration, lease details, staff information, growth opportunities, and owner's role. Unlocked only after the prospective buyer signs an NDA and provides proof of financial capacity. Morgan & Westfield
- **Stage 3 (Full Due Diligence, after LOI):** Tax returns, contracts, bank statements, employee records, legal documents, and complete operational data. OffDeal Access granted explicitly by the listing broker.

The PostgreSQL implementation uses a `listing_access_levels` table (user_id, listing_id, access_level, granted_by, granted_at) with middleware that checks: teaser → always allowed; CIM → requires valid NDA in `ndas` table; full DD → requires explicit grant plus valid NDA. For NDA signing at MVP, a simple "I agree" confirmation with timestamp and IP capture is sufficient. DocuSign's free tier (5 envelopes/month) or HelloSign can be added later.

The co-brokering workflow needs a `co_broke_agreements` table tracking listing broker, cooperating broker, commission split percentage, and acceptance status. When a cooperating broker's buyer signs an NDA, the cooperating broker automatically receives CIM access. The deal pipeline (`deal_activities` table) tracks the full lifecycle: inquiry → NDA signed → showing → offer → LOI → due diligence → closed.

Real estate MLS architecture provides the blueprint: standardized data schema, role-based access control, status tracking (active/pending/sold/withdrawn), and syndication feeds. The RESO Data Dictionary's approach of defining universal field names and values MLS API across 500+ MLSs should inform the broker listing schema MEV — standardize on fields like `industry_naics`, `revenue_range`, `ebitda_range`, `asking_price_range`, `employee_count_range`, `years_in_business`, and `reason_for_sale_category`.

The full stack costs $60–80/month at launch

The architecture runs entirely on Railway (Pro plan at $20/month), PostgreSQL for data storage

and caching, and a mix of free government APIs plus Google Places.

Cost category	Monthly estimate	Notes
Railway (Node.js + PostgreSQL + worker)	$20–30	Pro plan includes $20 credit; actual usage ~$21/month
Claude API (Haiku + Sonnet mix)	$15–40	Haiku at $1/$5 per MTok for bulk enrichment; Sonnet at $3/$15 for complex analysis
Google Places API	$0–50	Free within caps using field masks; Essentials IDs Only is unlimited
Census Bureau API	$0	Completely free with API key
SEC EDGAR	$0	Free, 10 req/sec, no authentication
SBA data (FOIA + lender reports)	$0	Free downloads, public domain
Transactional email	$0	SendGrid/Resend free tiers
Total	**$35–120**	Realistic average: ~$70/month

Claude API cost breakdown at 200 users: Business enrichment (2,000 profiles/month) costs $9–27 using Haiku; buyer matching (400 analyses/month) costs $3–12; deal reports (200/month) costs $2–8; ad-hoc queries cost $3–12. Prompt caching (sharing system prompts across requests) reduces input costs by up to **90%,** Claude API Docs and the Batch API provides an additional 50% discount Claude API Docs for non-time-sensitive enrichment jobs. Realistic monthly Claude spend: **$25–40**.

The caching strategy keeps API costs low. Since Google Places ToS prohibits permanent storage of detail data, Google the compliant approach caches Place IDs indefinitely Google while re-fetching details on demand. For non-Google data (Census, SBA, SEC, buyer profiles), PostgreSQL JSONB columns store API responses with TTLs: business listings at 30 days, PE firm profiles at 90 days, Census data at 365 days, search results at 24 hours. A cache-first pattern in Node.js checks PostgreSQL before making API calls, Redis and a nightly cleanup job purges expired entries. Materialized views aggregate cached data for common queries (businesses by region/industry) and refresh weekly. This eliminates **90%+ of redundant API calls**.

For data from multiple sources, deduplication keys on `normalized_name + zip_code` or `google_place_id` as the canonical identifier. A `data_sources` table tracks provenance: which

source provided which field, when, and with what confidence score. When sources conflict, prefer Google Places (freshest) → Census (official) → manual entry.

What ships in two to three weeks

Week 1: Foundation + acquisition target search. Import Census CBP data for the top 50 metros into PostgreSQL. Build a Google Places API integration with strict field masks (Essentials only). Implement the cache-first pattern. Create a search UI: user selects industry (NAICS) + geography → sees Census market density stats + live Google Places results with attribution. Import SBA 7(a) FOIA data and build a filter for businesses with acquisition-era loans.

Week 2: Buyer matching + professional directory. Import the SBIC directory CSV, Gigasheet PE/VC dataset, and manually curated search fund list into the buyer_profiles table. Build the overlap-scoring matching algorithm. When a seller enters their business profile, return top 10 matched buyers with Claude-generated match explanations. Simultaneously, manually curate 5–10 M&A professionals per category for the top 10 metros. Import SBA lender activity data. Build the simple professional directory UI with location/category filters.

Week 3: Broker hub MVP. Build the broker listing schema with teaser/CIM separation. Implement the NDA signing flow (simple "I agree" + timestamp). Build access-control middleware. Create the blind profile browse view and the CIM unlock flow. Add the co-brokering invitation workflow and commission split tracking. Build a basic deal pipeline tracker.

This delivers a platform where buyers can search for acquisition targets by industry and geography (powered by Census + Google Places + SBA data), sellers get matched to potential buyers (powered by SBIC + EDGAR + search fund data + AI scoring), users find M&A professionals (curated directory + SBA lender data), and brokers can list deals confidentially and collaborate — all running under $80/month in infrastructure costs. The approach mirrors Baton's strategy: **launch with public data, then build proprietary data as users engage.** Built In NYC Every seller who completes their profile, every buyer who registers criteria, and every broker who lists a deal makes the platform more valuable — creating the network effects that eventually compete with platforms spending orders of magnitude more on data acquisition.